

NNproject(BERT_model) (1)

April 20, 2023

```
[ ]: pip install transformers datasets evaluate
```

Looking in indexes: <https://pypi.org/simple>, <https://us-python.pkg.dev/colab-wheels/public/simple/>

Collecting transformers

Downloading transformers-4.28.1-py3-none-any.whl (7.0 MB)

7.0/7.0 MB

53.8 MB/s eta 0:00:00

Collecting datasets

Downloading datasets-2.11.0-py3-none-any.whl (468 kB)

468.7/468.7 kB

34.6 MB/s eta 0:00:00

Collecting evaluate

Downloading evaluate-0.4.0-py3-none-any.whl (81 kB)

81.4/81.4 kB

9.2 MB/s eta 0:00:00

Requirement already satisfied: regex!=2019.12.17 in

/usr/local/lib/python3.9/dist-packages (from transformers) (2022.10.31)

Requirement already satisfied: requests in /usr/local/lib/python3.9/dist-packages (from transformers) (2.27.1)

Collecting huggingface-hub<1.0,>=0.11.0

Downloading huggingface_hub-0.13.4-py3-none-any.whl (200 kB)

200.1/200.1 kB

18.4 MB/s eta 0:00:00

Requirement already satisfied: tqdm>=4.27 in

/usr/local/lib/python3.9/dist-packages (from transformers) (4.65.0)

Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.9/dist-packages (from transformers) (23.1)

Requirement already satisfied: numpy>=1.17 in /usr/local/lib/python3.9/dist-packages (from transformers) (1.22.4)

Requirement already satisfied: filelock in /usr/local/lib/python3.9/dist-packages (from transformers) (3.11.0)

Collecting tokenizers!=0.11.3,<0.14,>=0.11.1

Downloading

tokenizers-0.13.3-cp39-cp39-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (7.8 MB)

7.8/7.8 MB

```

88.0 MB/s eta 0:00:00
Requirement already satisfied: pyyaml>=5.1 in
/usr/local/lib/python3.9/dist-packages (from transformers) (6.0)
Requirement already satisfied: pandas in /usr/local/lib/python3.9/dist-packages
(from datasets) (1.5.3)
Collecting dill<0.3.7,>=0.3.0
  Downloading dill-0.3.6-py3-none-any.whl (110 kB)
      110.5/110.5 kB

11.6 MB/s eta 0:00:00
Collecting multiprocessing
  Downloading multiprocessing-0.70.14-py39-none-any.whl (132 kB)
      132.9/132.9 kB

15.0 MB/s eta 0:00:00
Requirement already satisfied: pyarrow>=8.0.0 in
/usr/local/lib/python3.9/dist-packages (from datasets) (9.0.0)
Collecting xxhash
  Downloading
xxhash-3.2.0-cp39-cp39-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (212 kB)
      212.2/212.2 kB

21.6 MB/s eta 0:00:00
Requirement already satisfied: fsspec[http]>=2021.11.1 in
/usr/local/lib/python3.9/dist-packages (from datasets) (2023.4.0)
Collecting aiohttp
  Downloading
aiohttp-3.8.4-cp39-cp39-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (1.0 MB)
      1.0/1.0 MB

56.1 MB/s eta 0:00:00
Collecting responses<0.19
  Downloading responses-0.18.0-py3-none-any.whl (38 kB)
Collecting async-timeout<5.0,>=4.0.0a3
  Downloading async_timeout-4.0.2-py3-none-any.whl (5.8 kB)
Collecting multidict<7.0,>=4.5
  Downloading
multidict-6.0.4-cp39-cp39-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (114
kB)
      114.2/114.2 kB

11.4 MB/s eta 0:00:00
Collecting yarl<2.0,>=1.0
  Downloading
yarl-1.8.2-cp39-cp39-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (264 kB)
      264.6/264.6 kB

23.6 MB/s eta 0:00:00
Collecting frozenlist>=1.1.1
  Downloading frozenlist-1.3.3-cp39-cp39-manylinux_2_5_x86_64.manylinux1_x86_64.
manylinux_2_17_x86_64.manylinux2014_x86_64.whl (158 kB)
      158.8/158.8 kB

15.8 MB/s eta 0:00:00
Collecting aiosignal>=1.1.2

```

Downloading aiosignal-1.3.1-py3-none-any.whl (7.6 kB)
 Requirement already satisfied: charset-normalizer<4.0,>=2.0 in
 /usr/local/lib/python3.9/dist-packages (from aiohttp->datasets) (2.0.12)
 Requirement already satisfied: attrs>=17.3.0 in /usr/local/lib/python3.9/dist-
 packages (from aiohttp->datasets) (23.1.0)
 Requirement already satisfied: typing-extensions>=3.7.4.3 in
 /usr/local/lib/python3.9/dist-packages (from huggingface-
 hub<1.0,>=0.11.0->transformers) (4.5.0)
 Requirement already satisfied: urllib3<1.27,>=1.21.1 in
 /usr/local/lib/python3.9/dist-packages (from requests->transformers) (1.26.15)
 Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.9/dist-
 packages (from requests->transformers) (3.4)
 Requirement already satisfied: certifi>=2017.4.17 in
 /usr/local/lib/python3.9/dist-packages (from requests->transformers) (2022.12.7)
 Requirement already satisfied: python-dateutil>=2.8.1 in
 /usr/local/lib/python3.9/dist-packages (from pandas->datasets) (2.8.2)
 Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.9/dist-
 packages (from pandas->datasets) (2022.7.1)
 Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.9/dist-
 packages (from python-dateutil>=2.8.1->pandas->datasets) (1.16.0)
 Installing collected packages: tokenizers, xxhash, multidict, frozenlist, dill,
 async-timeout, yarl, responses, multiprocessing, huggingface-hub, aiosignal,
 transformers, aiohttp, datasets, evaluate
 Successfully installed aiohttp-3.8.4 aiosignal-1.3.1 async-timeout-4.0.2
 datasets-2.11.0 dill-0.3.6 evaluate-0.4.0 frozenlist-1.3.3 huggingface-
 hub-0.13.4 multidict-6.0.4 multiprocessing-0.70.14 responses-0.18.0
 tokenizers-0.13.3 transformers-4.28.1 xxhash-3.2.0 yarl-1.8.2

```
[ ]: import torch
```

```
[ ]: from huggingface_hub import notebook_login

notebook_login()
```

Token is valid.
 Your token has been saved in your configured git credential helpers (store).
 Your token has been saved to /root/.cache/huggingface/token
 Login successful

```
[ ]: from datasets import load_dataset

squad = load_dataset("squad", split="train[:5000]")

squad = squad.train_test_split(test_size=0.2)

squad["train"][0]
{'answers': {'answer_start': [515], 'text': ['Saint Bernadette Soubirous']},
```

```

'context': 'Architecturally, the school has a Catholic character. Atop the
↳Main Building\'s gold dome is a golden statue of the Virgin Mary.
↳Immediately in front of the Main Building and facing it, is a copper statue
↳of Christ with arms upraised with the legend "Venite Ad Me Omnes". Next to
↳the Main Building is the Basilica of the Sacred Heart. Immediately behind
↳the basilica is the Grotto, a Marian place of prayer and reflection. It is a
↳replica of the grotto at Lourdes, France where the Virgin Mary reputedly
↳appeared to Saint Bernadette Soubirous in 1858. At the end of the main drive
↳(and in a direct line that connects through 3 statues and the Gold Dome), is
↳a simple, modern stone statue of Mary.',
'id': '5733be284776f41900661182',
'question': 'To whom did the Virgin Mary allegedly appear in 1858 in Lourdes
↳France?',
'title': 'University_of_Notre_Dame'
}

```

```

Downloading builder script: 0%|          | 0.00/5.27k [00:00<?, ?B/s]
Downloading metadata: 0%|          | 0.00/2.36k [00:00<?, ?B/s]
Downloading readme: 0%|          | 0.00/7.67k [00:00<?, ?B/s]
Downloading and preparing dataset squad/plain_text to /root/.cache/huggingface/d
atasets/squad/plain_text/1.0.0/d6ec3ceb99ca480ce37cdd35555d6cb2511d223b9150cce08
a837ef62ffea453...
Downloading data files: 0%|          | 0/2 [00:00<?, ?it/s]
Downloading data: 0%|          | 0.00/8.12M [00:00<?, ?B/s]
Downloading data: 0%|          | 0.00/1.05M [00:00<?, ?B/s]
Extracting data files: 0%|          | 0/2 [00:00<?, ?it/s]
Generating train split: 0%|          | 0/87599 [00:00<?, ? examples/s]
Generating validation split: 0%|          | 0/10570 [00:00<?, ? examples/s]
Dataset squad downloaded and prepared to /root/.cache/huggingface/datasets/squad
/plain_text/1.0.0/d6ec3ceb99ca480ce37cdd35555d6cb2511d223b9150cce08a837ef62ffea4
53. Subsequent calls will reuse this data.

```

```

[ ]: {'answers': {'answer_start': [515], 'text': ['Saint Bernadette Soubirous']},
      'context': 'Architecturally, the school has a Catholic character. Atop the Main
Building\'s gold dome is a golden statue of the Virgin Mary. Immediately in
front of the Main Building and facing it, is a copper statue of Christ with arms
upraised with the legend "Venite Ad Me Omnes". Next to the Main Building is the
Basilica of the Sacred Heart. Immediately behind the basilica is the Grotto, a
Marian place of prayer and reflection. It is a replica of the grotto at Lourdes,
France where the Virgin Mary reputedly appeared to Saint Bernadette Soubirous in
1858. At the end of the main drive (and in a direct line that connects through 3
statues and the Gold Dome), is a simple, modern stone statue of Mary.',

```

```

'id': '5733be284776f41900661182',
'question': 'To whom did the Virgin Mary allegedly appear in 1858 in Lourdes
France?',
'title': 'University_of_Notre_Dame'}

```

```

[ ]: from transformers import AutoTokenizer

tokenizer = AutoTokenizer.from_pretrained("bert-base-uncased")

def preprocess_function(examples):
    questions = [q.strip() for q in examples["question"]]
    inputs = tokenizer(
        questions,
        examples["context"],
        max_length=384,
        truncation="only_second",
        return_offsets_mapping=True,
        padding="max_length",
    )

    offset_mapping = inputs.pop("offset_mapping")
    answers = examples["answers"]
    start_positions = []
    end_positions = []

    for i, offset in enumerate(offset_mapping):
        answer = answers[i]
        start_char = answer["answer_start"][0]
        end_char = answer["answer_start"][0] + len(answer["text"][0])
        sequence_ids = inputs.sequence_ids(i)

        # Find the start and end of the context
        idx = 0
        while sequence_ids[idx] != 1:
            idx += 1
        context_start = idx
        while sequence_ids[idx] == 1:
            idx += 1
        context_end = idx - 1

        # If the answer is not fully inside the context, label it (0, 0)
        if offset[context_start][0] > end_char or offset[context_end][1] < start_char:
            start_positions.append(0)
            end_positions.append(0)
        else:
            # Otherwise it's the start and end token positions

```

```

idx = context_start
while idx <= context_end and offset[idx][0] <= start_char:
    idx += 1
start_positions.append(idx - 1)

idx = context_end
while idx >= context_start and offset[idx][1] >= end_char:
    idx -= 1
end_positions.append(idx + 1)

inputs["start_positions"] = start_positions
inputs["end_positions"] = end_positions
return inputs

```

```

Downloading (...)okenizer_config.json: 0%|          | 0.00/28.0 [00:00<?, ?B/s]
Downloading (...)lve/main/config.json: 0%|          | 0.00/570 [00:00<?, ?B/s]
Downloading (...)solve/main/vocab.txt: 0%|          | 0.00/232k [00:00<?, ?B/s]
Downloading (...) /main/tokenizer.json: 0%|          | 0.00/466k [00:00<?, ?B/s]

```

```

[ ]: tokenized_squad = squad.map(preprocess_function, batched=True,
    ↪remove_columns=squad["train"].column_names)

```

```

Map: 0%|          | 0/4000 [00:00<?, ? examples/s]
Map: 0%|          | 0/1000 [00:00<?, ? examples/s]

```

```

[ ]: from transformers import DefaultDataCollator

```

```

data_collator = DefaultDataCollator()

```

```

[ ]: from transformers import AutoModelForQuestionAnswering, TrainingArguments,
    ↪Trainer

```

```

model = AutoModelForQuestionAnswering.from_pretrained("bert-base-uncased")

```

```

training_args = TrainingArguments(
    output_dir="my_awesome_qa_model",
    evaluation_strategy="epoch",
    learning_rate=2e-5,
    per_device_train_batch_size=16,
    per_device_eval_batch_size=16,
    num_train_epochs=3,
    weight_decay=0.01,
    push_to_hub=True,
)

```

```

trainer = Trainer(
    model=model,
    args=training_args,
    train_dataset=tokenized_squad["train"],
    eval_dataset=tokenized_squad["test"],
    tokenizer=tokenizer,
    data_collator=data_collator,
)

trainer.train()

```

Downloading pytorch_model.bin: 0% | 0.00/440M [00:00<?, ?B/s]

Some weights of the model checkpoint at bert-base-uncased were not used when initializing BertForQuestionAnswering: ['cls.predictions.decoder.weight', 'cls.predictions.transform.LayerNorm.weight', 'cls.predictions.bias', 'cls.seq_relationship.bias', 'cls.predictions.transform.LayerNorm.bias', 'cls.seq_relationship.weight', 'cls.predictions.transform.dense.bias', 'cls.predictions.transform.dense.weight']

- This IS expected if you are initializing BertForQuestionAnswering from the checkpoint of a model trained on another task or with another architecture (e.g. initializing a BertForSequenceClassification model from a BertForPreTraining model).

- This IS NOT expected if you are initializing BertForQuestionAnswering from the checkpoint of a model that you expect to be exactly identical (initializing a BertForSequenceClassification model from a BertForSequenceClassification model). Some weights of BertForQuestionAnswering were not initialized from the model checkpoint at bert-base-uncased and are newly initialized: ['qa_outputs.bias', 'qa_outputs.weight']

You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.

Cloning https://huggingface.co/Rekhni/my_awesome_qa_model into local empty directory.

WARNING:huggingface_hub.repository:Cloning

https://huggingface.co/Rekhni/my_awesome_qa_model into local empty directory.

Download file pytorch_model.bin: 0% | 8.00k/253M [00:00<?, ?B/s]

Download file runs/Apr20_04-15-21_24ce5b36a57e/1681964126.999821/events.out.tfevents.1681964126.24ce5b36a57e.4...

Download file training_args.bin: 100%|#####| 3.50k/3.50k [00:00<?, ?B/s]

Clean file runs/Apr20_04-15-21_24ce5b36a57e/1681964126.999821/events.out.tfevents.1681964126.24ce5b36a57e.472....

Clean file training_args.bin: 29%|### | 1.00k/3.50k [00:00<?, ?B/s]

Download file runs/Apr20_04-15-21_24ce5b36a57e/events.out.tfevents.1681964126.24ce5b36a57e.472.0: 100%|#####...

Clean file runs/Apr20_04-15-21_24ce5b36a57e/events.out.tfevents.1681964126.

↳24ce5b36a57e.472.0: 20%|#9 ...

Clean file pytorch_model.bin: 0%| | 1.00k/253M [00:00<?, ?B/s]

/usr/local/lib/python3.9/dist-packages/transformers/optimization.py:391:

FutureWarning: This implementation of AdamW is deprecated and will be removed in a future version. Use the PyTorch implementation torch.optim.AdamW instead, or set `no_deprecation_warning=True` to disable this warning

warnings.warn(

<IPython.core.display.HTML object>

```
[ ]: TrainOutput(global_step=750, training_loss=1.8780564371744792,
metrics={'train_runtime': 27447.0103, 'train_samples_per_second': 0.437,
'train_steps_per_second': 0.027, 'total_flos': 2351670810624000.0, 'train_loss':
1.8780564371744792, 'epoch': 3.0})
```

```
[ ]: question = "How many programming languages does BLOOM support?"
context = "BLOOM has 176 billion parameters and can generate text in 46_
↳languages natural languages and 13 programming languages."
```

```
[ ]: from transformers import pipeline

question_answerer = pipeline("question-answering", model="my_awesome_qa_model")
question_answerer(question=question, context=context)
```

```
[ ]: {'score': 0.23441480100154877,
'start': 58,
'end': 95,
'answer': '46 languages natural languages and 13'}
```